


A RAJ MODI PRODUCTION

82 tests

sha256(prev + step)

ROC-AUC 0.913



★★★★★
"Exactly what we needed."
— Future Manager

★★★★★
"Stunning. Simply stunning."
— TBD Technical Interviewer

SELECT only

P = 0.0021

VERITA

Built from nothing. Built for you.


STARRING FASTAPI · SCIKIT-LEARN · DUCKDB · REACT
FEATURING CLAUDE CODE AS "THE AI MENTOR"
EXECUTIVE PRODUCER WOLTERS KLUWER (HOPEFULLY)
RUNNING TIME [X] WEEKS · RATED H FOR HONEST

Raj Modi · [INTERVIEW_DATE] · Wolters Kluwer Internship Interview


coming soon to a compliance department near you · no numbers were fabricated in the making of this product

THE ORIGIN STORY


A TRUE STORY · SIX PANELS · NO STUNT DOUBLES




THE PROBLEM




THE BRIEF




THE FIRST VERSION



CLAUDE CODE ARRIVES





THE BREAKTHROUGH



TODAY

A PROBLEM I NOTICED. A SOLUTION I BUILT. FOR YOU.



THE WORLD BEFORE

- An hour of BI setup before the first chart appears
- AI flags transactions and can't say *why*
- Metrics quoted without their evaluation discipline
- Agent "reasoning" that can be quietly rewritten
- Compliance teams asked to defend numbers they can't trace



THE WORLD AFTER

- One file drop → a BI-grade editable dashboard
- Every flag lists the exact signals that drove it
- Every metric is held-out and says so on screen
- Investigations sealed in a tamper-evident hash chain
- "How was this computed?" is a button, everywhere

82

TESTS, GREEN IN CI

11


FLAGSHIP FEATURES

0

FABRICATED METRICS

THE UNFAIR ADVANTAGE

OF NOT BEING A TRADITIONAL ENGINEER



I ask “why” before “how”

Verita exists because I interrogated an industry’s constraint — defensibility — before writing a line of code. The thesis came first; the stack served it.

I see problems before solutions

Reading your JD like a crime scene is the skill. The 500 generic AI portfolios this season prove how rare “what does this team actually fear?” still is.

I build for users, not reviewers

The “how was this computed?” button exists because an analyst will be *asked*. Features start at the person holding the consequence.

I explain technical things to anyone

This bundle has a grandmother-grade explanation and an engineer-grade one — of the same system, both accurate. That range is the JD’s “stakeholder communication,” demonstrated.

The argument, made boldly: pair this brain with your engineers and you don’t get a worse engineer — you get a translator who ships. The evidence is the rest of this bundle.

— THE CANDIDATE, UNHEDGED

ROUTE: ZERO → SHIPPED

ETA: [INTERVIEW_DATE]

fastest route, via honesty · tolls: one rebuilt flagship

time remaining: one interview

DEPARTURE: THE IDEA

read the JD, saw the gap

FIRST ENDPOINT

a file went in; JSON came out; heartbeat

AUTO-DASHBOARD CITY

profiler, recommender, editable canvas

IMBALANCE CANYON

99.8% accuracy, 0% usefulness — detour via PR-AUC

RISK ENGINE RIDGE

GBM + IsolationForest, honest 0.913

INJECTION PASS △

own tests drew blood; rebuilt the wall at engine level

CHAIN SEAL SUMMIT

SHA-256 links, tamper tests passing

THE INVESTIGATOR REBUILD

lost a week to honesty; found the flagship

HARDENING FLATS

82 tests, CI, Docker, 0.36s boot

ARRIVAL: THIS ROOM

bundle on the table, engine still warm

scenic overlook: the night it first worked

ROADTRIP PLAYLIST


- ▶ Now Playing: "Building Something Real"
- 2. Up Next: "Wolters Kluwer Internship"
- 3. Coming Soon: "V2 (feat. SHAP)"
- 4. Deep Cut: "0.36 Seconds to Boot"

TRIP COMPUTER

distance: zero → shipped FCC product in [X] weeks
fuel used: curiosity, stubbornness, one AI copilot
breakdowns: 2 major (both became war stories)
souvenirs: 82 tests, 11 flagship features, 1 thesis


WHAT I LEARNED — A SKY SURVEY

BRIGHT STARS: BRING ON DAY ONE · DIM STARS: HONEST, AND WORKING ON IT




THE NAVIGATOR

*Charts unknown technical territory.
Strongest when the path is unclear.*




THE ALCHEMIST

*Turns raw uploads into evidence.
Refuses to transmute by guesswork.*




THE GUARDIAN

*Earned its shield the hard way:
its own arrows, its own walls.*




THE MESSENGER

*Carries data across boundaries
without ever changing the message.*



THE ORACLE

*Speaks only with error bars attached.
Distrusts even its own prophecies.*



THE CARTOGRAPHER

*Maps everything, including the dragons.
Especially the dragons.*

*The bright ones I'll bring on day one.
The dim ones — I'll be working on by week two.*

magnitude scale: ● strong (shipped, tested) ● working (used, not mastered) ● developing (named, not faked)

I DID MY HOMEWORK

WHY WOLTERS KLUWER, SPECIFICALLY — WITH RECEIPTS

YOUR WORLD	MY UNDERSTANDING	MY CONTRIBUTION
WK sells <i>expert solutions</i> — software professionals bet their license on, in health, tax, legal, finance.	Your real product is trust under accountability; a wrong answer here has a legal consequence, not a refresh button.	Verita's spine is exactly that value: zero fabricated metrics, every claim traceable. I build like your customers audit.
FSS/FCC serves compliance teams fighting financial crime under BSA/AML/OFAC regimes.	Analysts drown in alerts; regulators demand that every flag be defensible; explainability is law, not luxury.	A fraud model with honest held-out metrics, an alert queue ranked by evidence, and an NLP analyzer that lists its signals.
Your JD asks for AI-driven dashboards, forecasting, NLP, GenAI — <i>plus</i> stakeholder communication.	The team wants builders whose AI work survives scrutiny and whose explanations survive non-engineers.	Every JD line maps to a shipped feature (the trace document is in the repo) — and this bundle is the communication, performed.
The industry's next decade: putting agentic AI in front of auditors without losing the trust the brand is built on.	Whoever solves AI <i>provenance</i> — what did the machine do and can we prove it — owns that decade.	A working prototype of that answer: the hash-chained, re-runnable Investigator. Small scale, right values, ready to learn yours.

**I didn't build a project that could work for any company.
I built one that works for yours. There's a difference.**

WHAT THIS PROJECT ACTUALLY IS, IN NUMBERS

ALL REAL. REAL BEATS IMPRESSIVE.

82

AUTOMATED TESTS IN CI

284,807

REAL TRANSACTIONS

0.913

HELD-OUT ROC-AUC

11

FLAGSHIP FEATURES

3

PRODUCT PILLARS

6

SQL TESTS PER
INVESTIGATION

0.36s

MODEL WARM-BOOT

[X]

WEEKS TO SHIPPED

0

FABRICATED METRICS




The number that matters most is the last one. Everything else on this page can be verified *because* of it: drop any CSV into the product and watch the numbers recompute from your data, live. The demo has no script to fall off.

Even where a number is small, it is real — and in the candidate's industry of choice, a real 0.913 outranks a fabricated 0.99 every day of the audit.

— THE HONESTY POLICY, APPLIED TO MARKETING

WHAT HAPPENS NEXT



Three specific wants — not “to learn and grow”:

1. My first code review from someone who has shipped FCC-grade systems — I’ve reviewed my own work as hard as I can alone; the next standard lives here.
2. Contact with production-scale, genuinely messy institutional data — the gap between public datasets and reality is the gap between credible and correct.
3. To watch how senior data scientists defend a model to a stakeholder or auditor — because that performance is the skill my whole project points toward.

**I DIDN'T COME HERE
WITH A CV.**

**I CAME HERE
WITH A SOLUTION.**

rajmodi262@gmail.com · github.com/rajmodi262/verita

[MY_LINKEDIN] · [MY_PHONE]

[OR →
GitHub]