

VERITA

Financial Crime & Compliance Intelligence — where every number shows its work

Built for **Wolters Kluwer** · Raj Modi · June 2026

0.913

ROC-AUC · HELD-OUT

0.65

PR-AUC · IMBALANCED


284,807

REAL TRANSACTIONS

ACT I · THE PROBLEM

Power BI's real cost isn't the charts — it's the *first hour* before you see one. And compliance teams legally *cannot act on a number they can't defend* to a regulator.

- 🚫 Hours lost before the first insight
- 🚫 Black-box dashboards nobody can **audit**
- 🚫 **Fraud** buried in 284,807-row files



"a number you can't defend"

ACT II · THE BUILD

1

Profile

Pandas · SciPy

2

Recommend

scipy.stats

3

Score

scikit-learn

4

Investigate

DuckDB

5

Seal

SHA-256

A real **scikit-learn** pipeline (**GradientBoosting** + **IsolationForest**) on the **ULB** fraud set; a sandboxed **DuckDB SQL** engine; an autonomous agent that tests five **AML** hypotheses and **hash-chains** every step it takes.



ACT III · THE RESULT

Drop a file → an **auditable**, BI-grade dashboard in seconds — every number shows its formula, every model decision its **SHAP** attribution, every investigation sealed.

- ✅ **Held-out** 0.913 / 0.65 — honest
- ✅ Every figure **reproducible**
- ✅ File → filed **SAR**, one workflow

GLASS BOX · every number shows its work

precision / recall



SHA-256 sealed



AI-NATIVE

builds with the model in the loop



PROBLEM-FIRST

starts from the regulator's need



SHIPS CODE

not promises — it runs

"Most AI analytics are black boxes. Verita makes every number defensible — then seals the proof."